

docs

Ramin Zarebidoky (LiterallyTheOne)

25 Oct 2025

Introduction

Keras

Keras is a high-level API for building and training **Deep Learning Models**. It is designed to be a stand-alone project. But with the help of **TensorFlow**, **PyTorch**, and **Jax**, It can run on top of different hardware (e.g., **CPU**, **GPU**).

The fascinating thing about **Keras** is that it is super easy to get started with. You can train and test a model with only a few lines of code. It is a perfect way to learn **Deep Learning** concepts by practically seeing their effects.

Google Colab

There are so many ways available to run a **Deep Learning** code. One of the fastest and easiest way that doesn't require any installation, is **Google Colab**. Google colab is a free could-based platform that is powered by jupyter notebook. All the packages that we want for this tutorial is already installed in **Google Colab**. Also, every code that we run in this tutorial can be run on this platform. So, I highly recommend you to start with **Google Colab**. After you have become more comfortable with the packages and concepts, switch to a local platform like your personal computer.

All the codes that we talk about in this tutorial is available in the **GitHub**. Each tutorial has a link to its respective code, which you can find it at the bottom of each page. To load and run the codes in **Google Colab**, you can follow these steps.

- Open Google colab
- From **files** select **Open Notebook**
- Go to the **GitHub** section
- Copy the **URL** of the code
- Select the **.ipynb** file that you want

Here is an example of loading this tutorial's code:

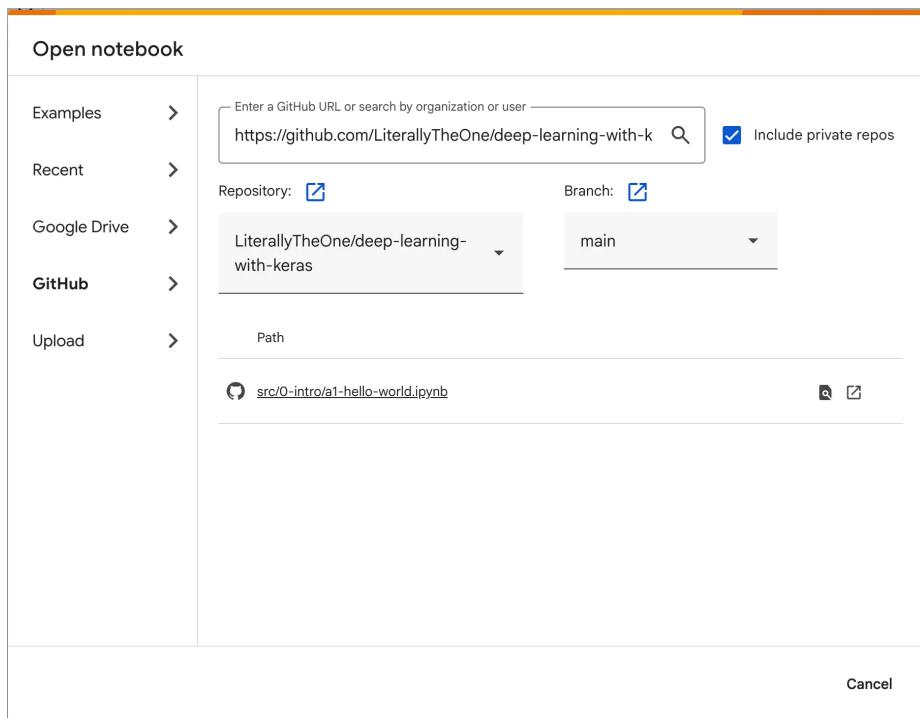


Figure 1: Colab GitHub

Hello World

Here is a **Hello World** example that we are gradually going to complete it step by step.

```
# Setup
import os

os.environ["KERAS_BACKEND"] = "torch"

# Imports
from keras.datasets import mnist
import keras
from keras import layers

# Prepare the Data
(train_images, train_labels), (test_images, test_labels) =
    mnist.load_data()

train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype("float32") / 255

# Define the model
model = keras.Sequential([
    layers.Dense(512, activation="relu"),
    layers.Dense(10, activation="softmax")
])

model.compile(optimizer="adam",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])

# Train the model
model.fit(train_images, train_labels, epochs=5, batch_size=128)

# Test the model
test_loss, test_acc = model.evaluate(test_images, test_labels)

"""

-----
output:

Epoch 1/5
469/469          2s 5ms/step - accuracy: 0.9259 - loss: 0.2622
Epoch 2/5
```

```

469/469          2s 5ms/step - accuracy: 0.9685 - loss: 0.1092
Epoch 3/5
469/469          2s 5ms/step - accuracy: 0.9797 - loss: 0.0710
Epoch 4/5
469/469          2s 5ms/step - accuracy: 0.9852 - loss: 0.0515
Epoch 5/5
469/469          2s 5ms/step - accuracy: 0.9901 - loss: 0.0363
313/313          1s 3ms/step - accuracy: 0.9801 - loss: 0.0616
"""

```

In the code above, we have trained and tested a model on a dataset called **MNIST**. In the future, we are going deeper into each step, but for now, here is a simple explanation of each of them. At first, we set up the backend of our **Keras**. We set that to **torch**, but you can set that to either **tensorflow** or **jax**. Then we imported the necessary modules. After that, we downloaded MNIST. MNIST contains of 28×28 images of handwritten digits between 0 and 9. Then, we normalize our data. After that, we defined a simple model and compiled the model with the proper **optimizer**, **loss**, and **metrics**. With the **fit** function, we train our model. And finally, we test our model with the **evaluate** function. As you can see in the output, our model's **accuracy** and **loss** are shown in each training step and testing step. We have gotten 99% accuracy on our training data and 98% accuracy on our test data.

Kaggle

Kaggle is one of the biggest platforms for data science and machine learning enthusiasts. It contains a huge number of datasets and a variety of competitions. In this tutorial, we are going to select an **Image Classification Dataset** from Kaggle. One of the simplest ways to do that is to go to the **Datasets** section in **Kaggle**, and select **Image Classification** tag in the **filters**. The dataset that we have to choose should have stored its images in a format like below:

```

class_a/
...a_image_1.jpg
...a_image_2.jpg
class_b/
...b_image_1.jpg
...b_image_2.jpg

```

As you can see, in the format above, we have some directories with images. Each directory represents a class, and the images in each directory belong to that class.

You can see the format of a **Dataset** by scrolling down to **Data Explorer**. For example, in Tom and Jerry Image classification We have a format as below:

As you can see, we have 4 directories (*jerry*, *tom*, *tom_jerry_0*, *tom_jerry_1*),

Data Explorer

Version 3 (469.3 MB)

- ▼ └ tom_and_jerry
 - ▼ └ tom_and_jerry
 - ▶ └ jerry
 - ▶ └ tom
 - ▶ └ tom_jerry_0
 - ▶ └ tom_jerry_1
- ████ challenges.csv
- ████ ground_truth.csv

Figure 2: Tom and Jerry data format

and each directory has its own images. So, we have 4 classes. Another example is Facial Emotion Recognition Dataset. Its data structure is as below:

Data Explorer

Version 1 (208.62 MB)

- ▼  processed_data
 - ▶  angry
 - ▶  disgust
 - ▶  fear
 - ▶  happy
 - ▶  neutral
 - ▶  sad
 - ▶  surprise

Figure 3: Facial Emotion data format

As you can see, in the image above, we have 7 directories (*angry*, *disgust*, *fear*, *happy*, *neutral*, *sad*, and *surprise*). So, we have 7 classes.

Now, you should select a dataset with these criteria:

- Image classification
- Each class has its own directory, and its images are in that directory
- It's better for our dataset size not to exceed 5GB.

Conclusion

In this tutorial, we have introduced **Keras**. Then, we explained about **Google Colab** and how to load a notebook from **GitHub**. After that, we provided a **Hello World** example that we are going to complete it overtime. Finally, we introduced **Kaggle** and explained how to get a suitable **Dataset** from it for this tutorial.

Load an Image Classification Dataset

Introduction

In the previous tutorial, we learned how about **Keras**, **Google Colab**, and **Kaggle**. Our task was to select an **Image Classification Dataset** from **Kaggle**. In this tutorial, we are going to load this dataset and make a ready to give it to a model.

Get data from Kaggle

The easiest and the recommended way to download a dataset from **Kaggle** is to use a package called **Kagglehub**. **Kaggle** itself has developed this package and made it super easy to use. You can learn more about this package in their GitHub Repository.

Now, how to use this package to download a dataset. In the dataset that you have selected, click on the **Download** button in the top right corner of the page. A window will pop up that has a code snippet on it. You should copy that code and use it in your own code. For Tom and Jerry Image classification, the is like this:

```
import kagglehub

# Download latest version
path = kagglehub.dataset_download(
    "balabaskar/tom-and-jerry-image-classification")

print("Path to dataset files:", path)
```

The code above, will automatically download the dataset and returns its path. We said that we wanted a structure like below:

```
class_a/
...a_image_1.jpg
...a_image_2.jpg
class_b/
...b_image_1.jpg
...b_image_2.jpg
```

We know that this dataset has this structure and if you looked at the dataset in **Kaggle**, you have noticed that it is in `tom_and_jerry/tom_and_jerry` directory. But get more familiar with the **jupyter notebook** commands, let's find it with taking the list of the path that we are currently on.

```
!ls {path}

"""
-----
output:

challenges.csv    ground_truth.csv tom_and_jerry
"""
```

As you can see, we have `tom_and_jerry` directory. Now, let's take the list of this directory.

```
!ls {path}/tom_and_jerry

"""
-----
output:

tom_and_jerry
"""
```

As you can see, we have another `tom_and_jerry` directory. Let's take the list of it to see what's inside of it.

```
!ls {path}/tom_and_jerry/tom_and_jerry

"""
-----
output:

jerry      tom      tom_jerry_0 tom_jerry_1
"""
```

And as you can see, we have reached to the structure that we wanted. Let's put this path in a variable called `data_path`, to be able to use it later.

```
from pathlib import Path

data_path = Path(path) / "tom_and_jerry/tom_and_jerry"
```

Your dataset might have subdirectories like `train`, `validation` and `test`. If it was like this put the `train` directory in the `data_path` and store the other ones in their respective directory. For example, `val_path` for validation and `test_path` for test.

ImageFolder

One of the best ways to use an **Image Classification Dataset** in PyTorch is by using `ImageFolder`. `ImageFolder` loads and assigns labels to a folder that has this structure:

```
main_directory/
...class_a/
.....a_image_1.jpg
.....a_image_2.jpg
...class_b/
.....b_image_1.jpg
.....b_image_2.jpg
```

This structure is the structure that we have right now in our `data_path` variable. Now, let's load our image folder and show one of the images.

```
from torchvision.datasets import ImageFolder
from matplotlib import pyplot as plt

all_data = ImageFolder(data_path)

for image, label in all_data:
    plt.figure()
    plt.imshow(image)
    print(label)
    break

"""
-----
output:
0
"""
```

As you can see, in the code above, we have loaded our images using `ImageFolder` and stored it in a variable called `all_data`. After that, we used a `for` to iterate through `images` and `labels`. We showed one image and one label and used `break` to end our loop. As it shown, the label is 0 and you can see the image representing that label in the above.

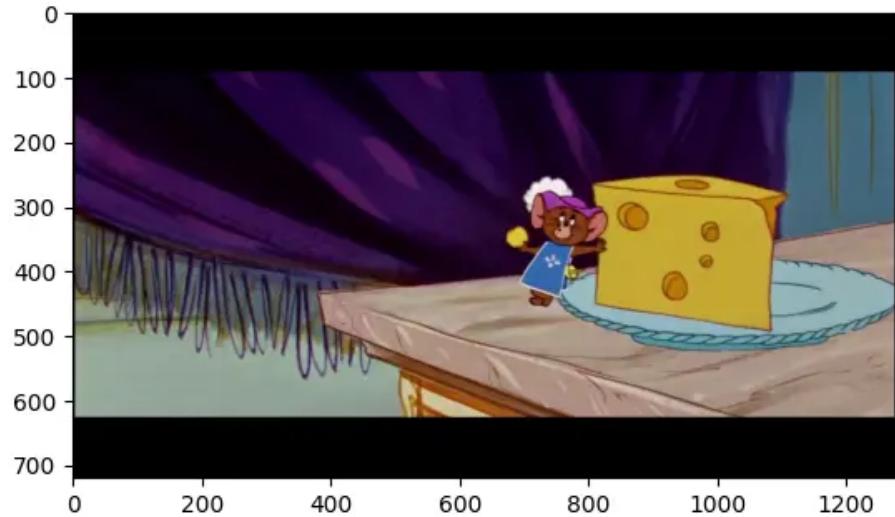


Figure 4: tom-and-jerry-example

Transforms

Transforms are the way that we can transform our images to the standard that we want. For example, when we load our dataset, the images might have different sizes. But when we want to train or test our model, we want images to have the same size. To make this happen, we can use the `transforms` module in `torchvision`. For example, let's load our dataset without a resize transform with `resize` transform and see the difference.

```
from torchvision import transforms

# Without resize transform

all_data = ImageFolder(data_path)

for image, label in all_data:
    print(f"image size without resize transform: {image.size}")
    break

# With resize transform

transform = transforms.Resize((90, 160))

all_data = ImageFolder(data_path, transform=transform)

for image, label in all_data:
```

```

    print(f"image size with resize transform: {image.size}")
    break

"""
-----
output:

image size without resize transform: (1280, 720)
image size with resize transform: (160, 90)
"""

```

As you can see, in the code above, we have successfully changed the size of our images to (160, 90).

Another thing is, when we load our images with `ImageFolder`, it would load them as PIL images. But when we want to feed our images to our model, we want them to be `tensors`. To achieve that, `torchvision` has a `transform` that takes an image and turns it into a `tensor`. To have resize and transforming to tensor transforms, we can combine them with each other like below:

```

trs = transforms.Compose(
    [
        transforms.Resize((90, 160)),
        transforms.ToTensor(),
    ]
)

all_data = ImageFolder(data_path, transform=trs)

for image, label in all_data:
    print(type(image))
    print(image.shape)
    break

"""
-----
output:

<class 'torch.Tensor'>
torch.Size([3, 90, 160])
"""

```

As you can see, we have our data in tensor, also the size of it is what we want it.

Split to train, validation, test

Some of our datasets don't have **train**, **validation**, **test** subsets. So, to split our data into these 3 subsets, we can use a function called **random_split**. This function, takes a **Dataset**, a sequence of **lengths** to split our data, and an optional **generator**. Here is an example on how to use **random_split**:

```
import torch
from torch.utils.data import random_split

g1 = torch.Generator().manual_seed(20)
train_data, val_data, test_data = random_split(all_data, [0.7,
    ↪ 0.2, 0.1], g1)

print(f"all_data's size: {len(all_data)}")
print(f"train_data's size: {len(train_data)}")
print(f"val_data's size: {len(val_data)}")
print(f"test_data's size: {len(test_data)}")

"""
-----
output:

all_data's size: 5478
train_data's size: 3835
val_data's size: 1096
test_data's size: 547
"""
```

In the code above, first we defined a **generator** with its seed set to 20. The reason for that is that we want every time that we run our code, have the same **train**, **validation**, and **test** subsets. Then, we used **random_split** function. for the first argument, we gave it **all_data** that we loaded it before. After that, we should give it a list of percentages or lengths. If we give it the percentages, sum of them should be equal to 1.0. If we give them the lengths, sum of them should be equal to the length of our data. For example [0.7, 0.2, 0.1] means to split data into 70, 20, and 10. We use that 70 for our training. We use 20 for validation. We use 10 for test. For the third argument, we gave the generator that we created earlier. As you can see in the result, we had 5478 samples, and we split them into train, validation, and test subsets. 3835 of them are for training, 1096 of them are for validation, and 547 of them are for testing.

For a **Deep Learning** project, we need these 3 subsets. If the **Dataset** provider hasn't split them already, we should split it. Otherwise, there is nothing to do.

DataLoader

Now, we have successfully loaded our dataset into tensors. Also, we have train, validation, and test subsets. Now, we are ready to feed them into our **model** for training and testing purposes. To make this procedure easier, **PyTorch** has a module called **DataLoader**. **Dataloader** takes a loaded dataset as its argument and helps us to apply the **Deep learning** techniques. One these techniques is called **mini-batch**. So, instead of feeding our data to our model one by one, we give it a **batch** of data. For example, each time we give it **12** data. It helps our model to learn better. Another technique is called **shuffling**. By **shuffling**, we change the order of data when we want to feed it to the model. It helps the model to learn more generally. To use **DataLoader** with these 2 techniques, we can use the code below:

```
train_loader = DataLoader(train_data, batch_size=12,
                           shuffle=True)
val_loader = DataLoader(val_data, batch_size=12, shuffle=False)
test_loader = DataLoader(test_data, batch_size=12, shuffle=False)
```

In the code above, we have 3 dataloaders for each train, validation, and test subsets. Then, we set the **batch_size** to 12 and for the train subset we set the **shuffle** to **true**. Now, let's show one batch of training data using **DataLoader**.

```
fig, axes = plt.subplots(3, 4)

axes_ravel = axes.ravel()

for images, labels in train_loader:
    for i, (image, label) in enumerate(zip(images, labels)):
        axes_ravel[i].imshow(transforms.ToPILImage()(image))
        axes_ravel[i].set_axis_off()
        axes_ravel[i].set_title(f"{label}")
    break
```

Output:

In the code above, we made a subplot with 3 rows and 4 columns. Then, we **ravel** it to make it a one dimensional array. This helps to use only one index instead of two. After that, we iterate thorough our **train_loader**. It would give us 12 images and 12 labels. Then we iterate through those images and labels and show them. As you recall, our images were in **tensor** format. To bring them back to PIL format, we can use a transform called **ToPILImage**. As you can see in the output, we have 12 different images with their respective label on top of them.



Figure 5: batch-tom-and-jerry

Your turn

Now, it is your turn. First, get your **Kaggle** dataset. Then, use the **ImageFolder** to load that dataset and show one of its images. After that, if you don't have any of the **train**, **validation**, and **test** subsets, make them using **random_split**. Then, load those three subsets using **DataLoader** and set a **batch_size** for them. Finally, show a batch of your data.

Conclusion

In this tutorial, we have learned how to work with a dataset. At first, we got an **Image classification** dataset from **Kaggle** using **Kagglehub**. Then, we loaded that dataset using **ImageFolder**. After that, we learned how to split our data if our dataset doesn't contain **train**, **validation**, and **test** subsets. Finally, we used **DataLoader** to load our data with **Deep Learning** techniques.