

Fine-tuning

Ramin Zarebidoky (LiterallyTheOne)

01 Sep 2025

Fine-tuning

Introduction

Fine-tuning is one of the most used techniques in **deep learning**. In this tutorial, we are going to learn how to load a pretrained model. Then, how to do **Transfer learning**. Finally, **Fine-tune** our model.

Load a dataset from Kaggle

In previous tutorials, we learned how to load a dataset from Kaggle. We have loaded a dataset called **Tom and Jerry image classification** and made the three subsets of **train**, **validation**, and **test**. Now, let's do it again.

```
path = kagglehub.dataset_download(_
    ↪ "balabaskar/tom-and-jerry-image-classification")
path = Path(path) / "tom_and_jerry/tom_and_jerry"

tom_and_jerry_transforms =
    ↪ transforms.Compose([transforms.Resize([90, 160]),
    ↪ transforms.ToTensor()])

all_data = ImageFolder(path, transform=tom_and_jerry_transforms)

g1 = torch.Generator().manual_seed(20)
train_data, val_data, test_data = random_split(all_data, [0.7,
    ↪ 0.2, 0.1], g1)

train_loader = DataLoader(train_data, batch_size=16,
    ↪ shuffle=True)
val_loader = DataLoader(val_data, batch_size=16, shuffle=False)
test_loader = DataLoader(test_data, batch_size=16, shuffle=False)
```

Let's plot on batch of its data:

```

images, labels = next(iter(train_loader))

fig, axes = plt.subplots(4, 4)

axes_ravel = axes.ravel()

for i, (image, label) in enumerate(zip(images, labels)):
    axes_ravel[i].imshow(transforms.ToPILImage()(image))
    axes_ravel[i].set_title(label.item())
    axes_ravel[i].set_axis_off()

```



Figure 1: Tom and Jerry Batch

Load a pretrained model

TorchVision has prepared some of the most famous vision models with pre-trained weights. In this tutorial, we are going to use a model called **MobileNetV2**. To load that model, we can use the code below:

```

from torchvision.models import mobilenet_v2, MobileNet_V2_Weights

model = mobilenet_v2(weights=MobileNet_V2_Weights.IMAGENET1K_V1)

```

```

print(model)

"""
-----
output:

MobileNetV2(
  (features): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2),
↪ padding=(1, 1), bias=False)
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
↪ track_running_stats=True)
      (2): ReLU6(inplace=True)
    )
    ...
    (17): InvertedResidual(
      (conv): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(160, 960, kernel_size=(1, 1), stride=(1,
↪ 1), bias=False)
          (1): BatchNorm2d(960, eps=1e-05, momentum=0.1,
↪ affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(960, 960, kernel_size=(3, 3), stride=(1,
↪ 1), padding=(1, 1), groups=960, bias=False)
          (1): BatchNorm2d(960, eps=1e-05, momentum=0.1,
↪ affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (2): Conv2d(960, 320, kernel_size=(1, 1), stride=(1, 1),
↪ bias=False)
        (3): BatchNorm2d(320, eps=1e-05, momentum=0.1,
↪ affine=True, track_running_stats=True)
      )
    )
    (18): Conv2dNormActivation(
      (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1),
↪ bias=False)
      (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1,
↪ affine=True, track_running_stats=True)
      (2): ReLU6(inplace=True)
    )
  )
)

```

```

(classifier): Sequential(
  (0): Dropout(p=0.2, inplace=False)
  (1): Linear(in_features=1280, out_features=1000, bias=True)
)
)

"""

```

In the code above, we have imported `mobilenet_v2` and `MobileNet_V2_Weights`. Then we loaded the mobile net with the weights that are pretrained on a dataset called **ImageNet**. Then, we printed the model to see the layers. As you can see, it consists of 2 **Sequential layers**. The first one is to extract features from the image. The second one is for classification.

Transfer Learning

Transfer learning is a technique of using a pretrained model (called the base model), on a new dataset with a different purpose. We don't train all the layers; instead, we only train the **classification layers**. So, the first step is to freeze all the layers.

```

# -----[ Freeze the model weights
↪ ]-----
for param in model.parameters():
    param.requires_grad = False

```

With the code above, we can freeze all the parameters. Now, let's replace the **classification layer** with our layer.

```

print("classifier before the change:")
print(model.classifier)
print("-" * 20)
# -----[ Change the classifier layer
↪ ]-----
model.classifier = nn.Linear(in_features=1280, out_features=4)

print("classifier after the change:")
print(model.classifier)

"""
-----
output:

classifier before the change:
Sequential(
  (0): Dropout(p=0.2, inplace=False)
  (1): Linear(in_features=1280, out_features=1000, bias=True)
)

```

```
)
-----
classifier after the change:
Linear(in_features=1280, out_features=4, bias=True)
"""
```

As you can see, in the code above, we have replaced the **classification layer** with our layer. As you recall, the dataset that we are using has 4 classes. So, the output of our final layer should be 4. Also, the output of the previous layer is 1280, so we should set our `in_features` to 1280 as well. Now let's see which layers are trainable:

```
for name, param in model.named_parameters():
    if param.requires_grad:
        print(f"{name} {param.shape}")

-----
output:

classifier.weight torch.Size([4, 1280])
classifier.bias torch.Size([4])
"""
```

In the code above, I have used `named_parameters` to go over the model parameters and also get their names. As you can see, the only parameters that are trainable are related to the **classifier**. So, let's replace **dataset** and the **model** that we had in `train_mnist_conv.py` and train our model. I have already applied the changes in `transfer_learning.py`. Now, let's run it for 20 epochs and see the results on **Tensorboard**.

```
-----
output:

mps
-----
epoch: 0
train:
    loss: 1.0890
    accuracy: 0.5330
validation:
    loss: 0.9879
    accuracy: 0.5894
```

```
-----
epoch: 1
train:
  loss: 0.9273
  accuracy: 0.6248
validation:
  loss: 0.8319
  accuracy: 0.6542
-----
...
-----
epoch: 19
train:
  loss: 0.6706
  accuracy: 0.7379
validation:
  loss: 0.6769
  accuracy: 0.7199
-----
test:
  loss: 0.7269
  accuracy: 0.7130
"""
```

Transfer Learning Train Accuracy




Figure 2: Transfer Learning Accuracy Train




Figure 3: Transfer Learning Accuracy Validation

Transfer Learning Validation Accuracy

As you can see, in the results above, we have reached acceptable accuracies.

subset	accuracy
train	73.79
validation	71.99
test	71.30

The results are pretty close, so our model is not overfitting. Also, the charts are ascending.

Fine-tuning

Fine-tuning has the same purpose as **Transfer Learning**. The only exception is that we train more layers. So, let's load our model again and freeze all layers except the last two (17 and 18).

```
model = mobilenet_v2(weights=MobileNet_V2_Weights.IMAGENET1K_V1)

# -----[ Freeze the model weights
↪ ]-----
for name, param in model.named_parameters():
    if not ("18" in name or "17" in name):
        param.requires_grad = False
```

In the code above, I iterated over the model's parameters. If they had 18 or 17 in their names, I didn't freeze them. Now, let's change the classifier layer and print the trainable parameters.

```

model.classifier = nn.Linear(in_features=1280, out_features=4)

for name, param in model.named_parameters():
    if param.requires_grad:
        print(f"{name} {param.shape}")

"""
-----
output:

features.17.conv.0.0.weight torch.Size([960, 160, 1, 1])
features.17.conv.0.1.weight torch.Size([960])
features.17.conv.0.1.bias torch.Size([960])
features.17.conv.1.0.weight torch.Size([960, 1, 3, 3])
features.17.conv.1.1.weight torch.Size([960])
features.17.conv.1.1.bias torch.Size([960])
features.17.conv.2.weight torch.Size([320, 960, 1, 1])
features.17.conv.3.weight torch.Size([320])
features.17.conv.3.bias torch.Size([320])
features.18.0.weight torch.Size([1280, 320, 1, 1])
features.18.1.weight torch.Size([1280])
features.18.1.bias torch.Size([1280])
classifier.weight torch.Size([4, 1280])
classifier.bias torch.Size([4])
"""

```

As you can see, the last two layers of our model are still trainable, and we have a classifier that works with our dataset. I have already applied the required changes in `fine_tuning.py`. Let's run it to see the results.

```

"""
-----
output:

mps
-----
epoch: 0
train:
    loss: 0.9683
    accuracy: 0.6188
validation:
    loss: 0.7647
    accuracy: 0.6870
-----

```



```

epoch: 1
train:
    loss: 0.6625
    accuracy: 0.7458
validation:
    loss: 0.5958
    accuracy: 0.7737
-----
...
-----
epoch: 18
train:
    loss: 0.1789
    accuracy: 0.9335
validation:
    loss: 0.5616
    accuracy: 0.8650
-----
epoch: 19
train:
    loss: 0.1397
    accuracy: 0.9518
validation:
    loss: 0.6718
    accuracy: 0.8577
-----
test:
    loss: 0.6284
    accuracy: 0.8537

"""

```

Fine-tuning Train Accuracy

- Orange: Transfer Learning
- Red: Fine-tuning

Fine-tuning Validation Accuracy

- Orange: Transfer Learning
- Red: Fine-tuning

As you can see in the results above, we have achieved better results than **Transfer Learning**.




Figure 4: Fine-tuning Train accuracy




Figure 5: Fine-tuning Validation Accuracy

subset	accuracy
train	95.18
validation	85.77
test	85.37

Conclusion

In this tutorial, we learned how to use a pretrained model on a new dataset. This is one of the most used techniques in deep learning. At first, we learned about **Transfer Learning** and saw the results. Then, we learned about **Fine-tuning** and compared it with **Transfer Learning**.